

# TEXT MINING

## WEB INFORMATION RETRIEVAL

Fang Wei-Kleiner

ADIT/IDA  
Linköping University

# OUTLINE

1. PageRank
2. Topic-Specific PageRank
3. Link Spam
4. A simple crawler

# EARLY WEB SEARCH

- How to organize the Web?
  - ▶ First try: Human curated Web directories Yahoo, DMOZ.
  - ▶ Second try: Web Search
    - Information Retrieval investigates:  
Find relevant docs in a small and trusted set
      - Newspaper
      - articles, Patents
- But: Web is huge, full of untrusted documents, random things, web spam, etc.

# EARLY WEB SEARCH ENGINE

- Early Web search engine worked by crawling the Web → terms in inverted index → query
- Ranked query processing:
  - ▶ Presence of a term in a header → higher rank
  - ▶ Large numbers of occurrences of the term → higher rank
- Term Spam

# TERM SPAM

- A T-shirt seller could add a term MOVIE to his page, and **do it thousands of times**.
- When a user issued a search query with the term MOVIE, the search engine would list that page first.
- Many tricks:
  - ▶ Give it the same color as the background.
  - ▶ Go to the search engine, issue the query MOVIE → copy the 1st ranked page → using the background color to make it invisible.
- **Term Spam**: techniques for fooling search engines into believing your page is about something it is not.
- Term spam rendered early search engines almost useless.

- PageRank was used to simulate where Web surfers
  - ▶ Starting at a random page
  - ▶ Would tend to congregate if they followed randomly chosen outlinks from the page at which they were currently located
  - ▶ This process were allowed to iterate many times.
  - ▶ Pages that would have a large number of surfers were considered more **important** than pages that would rarely be visited.
- Google prefers important pages to unimportant pages.
- Page judged not only by the terms appearing on that page, but by the terms used in or near the links to that page.
  - ▶ Spammer cannot easily get false terms added to these pages.


Ok, but why **simulation of random surfers** should allow us to approximate the intuitive notion of the **importance** of pages?

- Users of the Web **vote with their feet**.
  - They tend to place links to pages they think are good or useful pages to look at, rather than bad or useless pages.
- The behavior of a random surfer indicates which pages users of the Web are likely to visit.
  - Users are more likely to visit useful pages than useless pages.

**PageRank measure has been proved empirically to work.**

# PAGERANK: TRANSITION MATRIX

A hypothetical example of the Web



Transition matrix

$$M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{pmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{pmatrix} \end{matrix}$$

Element  $m_{ij}$  in row  $i$  and column  $j$  has value  $1/k$  if page  $j$  has  $k$  arcs out, and one of them is to page  $i$ . Otherwise,  $m_{ij} = 0$ .

- Model the Web as a directed graph. Pages: nodes, Links: edges.
- The **transition matrix** of the Web  $M$  has  $n$  rows and columns for the Web with  $n$  pages.



# PAGERANK: DEFINITION

## Definition (PageRank)

The probability distribution for the location of a random surfer can be described by a column vector whose  $j$ th component is the probability that the surfer is at page  $j$ . This probability is the (idealized) **PageRank** function.

- A random surfer at any of the  $n$  pages of the Web with equal probability. Then the initial vector  $v_0$  will have  $1/n$  for each component.
- If  $M$  is the transition matrix of the Web, then after **one** step, the distribution of the surfer will be  $Mv_0$ , after **two** steps it will be  $M(Mv_0) = M^2v_0 \dots$

→  $M^i v_0$  is the distribution of the surfer after  $i$  steps.

## PAGERANK: TRANSITION MATRIX

$$\begin{pmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{pmatrix} \times \begin{pmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{pmatrix} = \begin{pmatrix} 9/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{pmatrix}$$

$$\begin{pmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{pmatrix} \times \begin{pmatrix} 9/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{pmatrix} = \begin{pmatrix} 15/48 \\ 11/48 \\ 11/48 \\ 11/48 \end{pmatrix}$$

...

$$\begin{pmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{pmatrix} \times \begin{pmatrix} 3/9 \\ 2/9 \\ 2/9 \\ 2/9 \end{pmatrix} = \begin{pmatrix} 3/9 \\ 2/9 \\ 2/9 \\ 2/9 \end{pmatrix}$$

## PAGERANK: DEFINITION

$$\begin{pmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{pmatrix} \times \begin{pmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{pmatrix} = \begin{pmatrix} 9/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{pmatrix}$$

The probability  $x_i$  that a random surfer will be at node  $i$  at the next step, is

$$\sum_j m_{ij} v_j$$

where  $m_{ij}$  is the probability that a surfer at node  $j$  will move to node  $i$  at the next step and  $v_j$  is the probability that the surfer was at node  $j$  at the previous step.


- This behavior is an example of the theory of [Markov processes](#).

# PAGERANK: MARKOV PROCESS

- It is known that the distribution of the surfer approaches a limiting distribution  $v$  that satisfies  $v = Mv$ , provided two conditions are met:
  - ▶ The graph is strongly connected; that is, it is possible to get from any node to any other node.
  - ▶ There are no dead ends: nodes that have no arcs out.
- Limit reached means the limiting  $v$  is an **eigenvector** of  $M \rightarrow Mv = v$ .
- $M$  is **stochastic**  $\rightarrow$  its columns each add up to 1.
- The principal eigenvector of  $M$  tells us where the surfer is most likely to be after a long time.
- We can compute the principal eigenvector of  $M$  by starting with the initial vector  $v_0$  and multiplying by  $M$  some number of times, until the vector we get shows little change at each round.

# WEB PICTURE


## The bowtie picture of the Web



- In-component: could reach SCC, but not reachable from the SCC.
- Out-component: reachable from the SCC but unable to reach the SCC.
- Tendrils:
  - ▶ out: reachable from the in-component but not able to reach the in-component.
  - ▶ in: can reach out-component, but are not reachable from out-component.
- Tubes, isolated components

# WEB PICTURE

## The bowtie picture of the Web




### Problems:

- Violation on assumptions needed for the Markov process iteration to converge to a limit.
- Out-components: spider traps.
- Surfers starting at SCC, in-components eventually wind up in out-components or tendrils.
- Page in the SCC or in-component winds up with probability of 0.

# PAGERANK: DEAD END

With dead ends, the transition matrix of the Web is no longer stochastic  
→ some of the columns will sum to 0 rather than 1.

## Web with dead end



## Transition matrix

$$M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{pmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{pmatrix} \end{matrix}$$

C is a dead end. In terms of random surfers, when surfers reaches C they disappear at the next round.

## PAGERANK: DEAD END

Starting with the vector with each component  $1/4$ , and repeatedly multiplying the vector by  $M$ :

$$\begin{pmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{pmatrix} \begin{pmatrix} 3/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{pmatrix} \begin{pmatrix} 5/48 \\ 7/48 \\ 7/48 \\ 7/48 \end{pmatrix} \cdots \begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$


→ After some time, all the surfers will be landing on  $C$  and **drains out** of the Web.



# PAGERANK: DEAD END

With dead ends, the transition matrix of the Web is no longer stochastic  
→ some of the columns will sum to 0 rather than 1.

## Web with dead end



## Transition matrix

$$M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{pmatrix} 0 & 1/2 & \textcolor{red}{1/4} & 0 \\ 1/3 & 0 & \textcolor{red}{1/4} & 1/2 \\ 1/3 & 0 & \textcolor{red}{1/4} & 1/2 \\ 1/3 & 1/2 & \textcolor{red}{1/4} & 0 \end{pmatrix} \end{matrix}$$

Modify the process by simulating random surfers moving about the Web.

## PAGERANK: MODIFY PROCESS FOR DEAD END

Starting with the vector with each component  $1/4$ , and repeatedly multiplying the vector by  $M$ :


$$\begin{pmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{pmatrix} \begin{pmatrix} 9/48 \\ 13/48 \\ 13/48 \\ 13/48 \end{pmatrix} \begin{pmatrix} 39/192 \\ 51/192 \\ 51/192 \\ 51/192 \end{pmatrix} \begin{pmatrix} 153/768 \\ 205/768 \\ 205/768 \\ 205/768 \end{pmatrix} \cdots \begin{pmatrix} 3/15 \\ 4/15 \\ 4/15 \\ 4/15 \end{pmatrix}$$

→ Converges!

# PAGERANK: SPIDER TRAPS

A spider trap is a set of nodes with no dead ends but no arcs out.

## Web with spider traps



## Transition matrix

$$M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{pmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 1 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{pmatrix} \end{matrix}$$

C a simple spider trap of one node.  
Note that in general spider traps can have many nodes.

## PAGERANK: SPIDER TRAPS

Starting with the vector with each component  $1/4$ , and repeatedly multiplying the vector by  $M$ :


$$\begin{pmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{pmatrix} \begin{pmatrix} 3/24 \\ 5/24 \\ 11/24 \\ 5/24 \end{pmatrix} \begin{pmatrix} 5/48 \\ 7/48 \\ 29/48 \\ 7/48 \end{pmatrix} \begin{pmatrix} 21/288 \\ 31/288 \\ 205/288 \\ 31/288 \end{pmatrix} \cdots \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \end{pmatrix}$$

→ All the PageRank is at  $C$ , since once there a random surfer there, he can never leave.

# PAGERANK: APERIODIC GRAPHS

**Aperiodicity.** Roughly: The pages cannot be partitioned such that the random walker visits the partitions sequentially

Graph which is not aperiodic



Transition matrix

$$M = \begin{matrix} & \begin{matrix} A & B \end{matrix} \\ \begin{matrix} A \\ B \end{matrix} & \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix} \end{matrix}$$

Starting with the vector  $\begin{pmatrix} 1 \\ 0 \end{pmatrix}$ , and repeatedly multiplying the vector by  $M$ :

$$\begin{pmatrix} 1 \\ 0 \end{pmatrix} \begin{pmatrix} 0 \\ 1 \end{pmatrix} \begin{pmatrix} 1 \\ 0 \end{pmatrix} \cdots$$

# ERGODIC MARKOV CHAINS

- A Markov chain is ergodic iff it is irreducible and aperiodic.
- **Irreducibility.** Roughly: there is a path from any page to any other page.
- **Aperiodicity.** Roughly: The pages cannot be partitioned such that the random walker visits the partitions sequentially.

## PAGERANK: 3 QUESTIONS

$$Mv = v$$

- Does this converge?  
→ no. As long as the graph does not fulfill those conditions. Modifying the graphs is not a good idea.
- Does it converge to what we want?  
→ no. It does not really describe the random surfer's behaviour.
- Are results reasonable?  
→ no. A surfer does not simply stop or get trapped repeatedly. She can always jump out and start a new page.

# PAGERANK: TELEPORTING

- We modify the calculation of PageRank by allowing each random surfer a small probability of teleporting to a random page, rather than following an out-link from their current page.
- The iterative step, where we compute a new vector estimate of PageRanks  $v'$  from the current PageRank estimate  $v$  and the transition matrix  $M$  is

$$v' = \beta Mv + (1 - \beta)e/n$$



# PAGERANK: TELEPORTING

$$v' = \beta Mv + (1 - \beta)e/n$$

- $\beta$ : a chosen constant, usually in the range 0.8 to 0.9.
- $e$ : a vector of all 1's with the appropriate number of components.
- $n$ : the number of nodes in the Web graph.
- $\beta Mv$  represents the case where, with probability  $\beta$ , the random surfer decides to follow an out-link from their present page.
- The term  $(1 - \beta)e/n$  is a vector each of whose components has value  $(1 - \beta)/n$  and represents the introduction, with probability  $(1 - \beta)$ , of a new random surfer at a random page.

## PAGERANK: TELEPORTING

$$\text{Let } M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{pmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 1 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{pmatrix} \end{matrix}$$

If we set  $\beta$  as 0.8, the equation for the iteration becomes

$$v' = \begin{pmatrix} 0 & 2/5 & 0 & 0 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 0 & 4/5 & 2/5 \\ 4/15 & 2/5 & 0 & 0 \end{pmatrix} v + \begin{pmatrix} 1/20 \\ 1/20 \\ 1/20 \\ 1/20 \end{pmatrix}$$

→ incorporated the factor  $\beta$  into  $M$  by multiplying each of its elements by  $4/5$ .

## PAGERANK: TELEPORTING

Starting with the vector with each component  $1/4$ , and repeatedly multiplying the vector by  $M$ :

$$\begin{pmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{pmatrix} \begin{pmatrix} 9/60 \\ 13/60 \\ 25/60 \\ 13/60 \end{pmatrix} \begin{pmatrix} 41/300 \\ 53/300 \\ 153/300 \\ 53/300 \end{pmatrix} \begin{pmatrix} 543/4500 \\ 707/4500 \\ 2543/4500 \\ 707/4500 \end{pmatrix} \cdots \begin{pmatrix} 15/148 \\ 19/148 \\ 95/148 \\ 19/148 \end{pmatrix}$$

→ By being a spider trap,  $C$  has managed to get more than half of the PageRank for itself.

# ERGODIC MARKOV CHAINS

- Theorem: For any ergodic Markov chain, there is a unique long-term visit rate for each state.
- This is the **steady-state probability distribution**.
- Over a long time period, we visit each state in proportion to this rate.
- It doesn't matter where we start.
- **Teleporting makes the process ergodic.**
- **$\Rightarrow$  Web-graph+teleporting has a steady-state probability distribution.**
- **$\Rightarrow$  Each page in the web-graph+teleporting has a PageRank.**

# TOPIC-SPECIFIC PAGERANK

- Instead of generic popularity, can we measure popularity within a topic?
- Goal: Evaluate Web pages not just according to their popularity, but by how close they are to a particular topic, e.g **sports** or **history**
- Allows search queries to be answered based on interests of the user
- Example: Query **Trojan** wants different pages depending on whether you are interested in sports, history and computer security

# TOPIC-SPECIFIC PAGERANK

- Random walker has a small probability of teleporting at any step
- Teleport can go to:
  - ▶ Standard PageRank: Any page with equal probability (To avoid dead end and spider trap problems)
  - ▶ Topic Specific PageRank: A topic specific set of **relevant** pages (teleport set)
- Idea: Bias the random walk
  - ▶ When walker teleports, she picks a page from a set  $S$
  - ▶  $S$  contains only pages that are relevant to the topic. → E.g., Open Directory (DMOZ) pages for a given topic/query
  - ▶ For each teleport set  $S$ , we get a different vector  $r_S$

# TOPIC-SPECIFIC PAGERANK


- Suppose  $S$  is a set of integers consisting of the numbers for the pages we have identified as belonging to a certain topic (called the **teleport set**).
- Let  $e_S$  be a vector that has 1 in the components in  $S$  and 0 in other components. Then the topic-specific PageRank for  $S$  is the limit of the iteration

$$v' = \beta Mv + (1 - \beta)e_S/|S|$$

where  $M$  is the transition matrix of the Web, and  $|S|$  is the size of set  $S$ .

# TOPIC-SPECIFIC PAGERANK

## A hypothetical example of the Web



## Transition matrix

$$M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{pmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{pmatrix} \end{matrix}$$

$$\beta M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{pmatrix} 0 & 2/5 & 4/5 & 0 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 2/5 & 0 & 0 \end{pmatrix} \end{matrix}$$

Where  $\beta = 0.8$ .



## TOPIC SPECIFIC PAGERANK

Suppose our topic is represented by the teleport set  $S = \{B, D\}$ . Then the vector  $(1 - \beta)e_S/|S|$  has  $1/10$  for its second and fourth components and 0 for the other two components. ( $1/10$  comes from  $0.2*1/2$ ).

$$v' = \begin{pmatrix} 0 & 2/5 & 4/5 & 0 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 2/5 & 0 & 0 \end{pmatrix} v + \begin{pmatrix} 0 \\ 1/10 \\ 0 \\ 1/10 \end{pmatrix}$$

$$\begin{pmatrix} 0/2 \\ 1/2 \\ 0/2 \\ 1/2 \end{pmatrix} \quad \begin{pmatrix} 2/10 \\ 3/10 \\ 2/10 \\ 3/10 \end{pmatrix} \quad \begin{pmatrix} 42/150 \\ 41/150 \\ 26/150 \\ 41/150 \end{pmatrix} \quad \begin{pmatrix} 62/250 \\ 71/250 \\ 46/250 \\ 71/250 \end{pmatrix} \quad \cdots \quad \begin{pmatrix} 54/210 \\ 59/210 \\ 38/210 \\ 59/210 \end{pmatrix}$$


→  $B$  and  $D$  get a higher PageRank than they did before.

# LINK SPAM

- Once Google became the dominant search engine, spammers began to work out ways to fool Google
- Spam farms were developed to concentrate PageRank on a single page
- Link spam: Creating link structures that boost PageRank of a particular page

# LINK SPAM

## The Web from the point of view of the link spammer



- A collection of pages whose purpose is to increase the PageRank of a certain page or pages is called a **spam farm**.
- **target page**  $t$ : at which spammer attempts to place as much PageRank as possible.
- A large number  $m$  of **supporting pages**: accumulate the portion of the PageRank that is distributed equally to all pages.

# ANALYSIS OF A SPAM FARM

- Taxation parameter  $\beta$ , typically around 0.85.
- $n$  be pages on the Web,  $m$  be the number of supporting pages.
- $x$  be the amount of PageRank contributed by the accessible pages.
  - ▶  $\rightarrow x$  is the sum, over all accessible pages  $p$  with a link to  $t$ , of the PageRank of  $p$  times  $\beta$ , divided by the number of successors of  $p$ .
- Let  $y$  be the unknown PageRank of  $t$ . We shall solve for  $y$ .

PageRank of each supporting page is  $\beta y/m + (1 - \beta)/n$

Then,

$$\begin{aligned}y &= x + \beta m(\beta y/m + (1 - \beta)/n) + (1 - \beta)/n(\text{ignored}) \\&= x/(1 - \beta^2) + c(m/n)\end{aligned}$$

where  $c = \beta/(1 + \beta)$ .

For  $\beta = 0.85$ ,  $(1 - \beta^2) = 3.6 \rightarrow$  amplified the external PageRank contribution by 360%. Increasing  $m$  will increase  $y$ .

# COMBATING LINK SPAM: TRUSTRANK

- **TrustRank**: topic specific PageRank with a teleport set of **trusted** pages. → Example: edu domains, similar domains for non US schools.
- Basic principle: while a spam page might easily be made to link to a trustworthy page, it is unlikely that a trustworthy page would link to a spam page.
- The borderline area is a site with blogs or other opportunities for spammers to create links. These pages cannot be considered trustworthy.

→ It is likely that search engines today implement this strategy routinely, so that what we think of as PageRank really is a form of TrustRank.

# HOW HARD CAN CRAWLING BE?

- Web search engines must crawl their documents.
- Getting the content of the documents is easier for many other IR systems.
  - ▶ E.g., indexing all files on your hard disk: just do a recursive descent on your file system
- Ok: for web IR, getting the content of the documents takes longer  
...
- ... because of latency.
- But is that really a design/systems challenge?

# BASIC CRAWLER OPERATION

- Initialize queue with URLs of known seed pages
- Repeat
  - ▶ Take URL from queue
  - ▶ Fetch and parse page
  - ▶ Extract URLs from page
  - ▶ Add URLs to queue
- Fundamental assumption: The web is well linked.

## EXERCISE: WHAT'S WRONG WITH THIS CRAWLER?

```
urlqueue := (some carefully selected set of seed urls)
while urlqueue is not empty:
    myurl := urlqueue.getlastanddelete()
    mypage := myurl.fetch()
    fetchedurls.add(myurl)
    newurls := mypage.extracturls()
    for myurl in newurls:
        if myurl not in fetchedurls and not in urlqueue:
            urlqueue.add(myurl)
    addtoinvertedindex(mypage)
```



# WHAT'S WRONG WITH THE SIMPLE CRAWLER

- Scale: we need to **distribute**.
- We can't index everything: we need to **subselect**. How?
- Duplicates: need to integrate **duplicate detection**
- Spam and spider traps: need to integrate **spam detection**
- **Politeness**: we need to be “nice” and space out all requests for a site over a longer period (hours, days)
- **Freshness**: we need to recrawl periodically.
  - ▶ Because of the size of the web, we can do frequent recrawls only for a small subset.
  - ▶ Again, subselection problem or **prioritization**

# MAGNITUDE OF THE CRAWLING PROBLEM

- To fetch 20,000,000,000 pages in one month . . .
- . . . we need to fetch almost 8000 pages per second!
- Actually: many more since many of the pages we attempt to crawl will be duplicates, unfetchable, spam etc.

# WHAT A CRAWLER MUST DO

## Be polite

- Don't hit a site too often
- Only crawl pages you are allowed to crawl: robots.txt

## Be robust

- Be immune to spider traps, duplicates, very large pages, very large websites, dynamic pages etc

# ROBOTS.TXT

- Protocol for giving crawlers (“robots”) limited access to a website, originally from 1994
- Examples:
  - ▶ User-agent: \*  
Disallow: /yoursite/temp/
  - ▶ User-agent: searchengine  
Disallow: /
- Important: cache the robots.txt file of each site we are crawling

## EXAMPLE OF A ROBOTS.TXT (NIH.GOV)

```
User-agent: PicoSearch/1.0
Disallow: /news/information/knight/
Disallow: /nidcd/
...
Disallow: /news/research_matters/secure/
Disallow: /od/ocpl/wag/
User-agent: *
Disallow: /news/information/knight/
Disallow: /nidcd/
...
Disallow: /news/research_matters/secure/
Disallow: /od/ocpl/wag/
Disallow: /ddir/
Disallow: /sdminutes/
```

# WHAT ANY CRAWLER SHOULD DO

- Be capable of **distributed** operation
- Be scalable: need to be able to increase crawl rate by adding more machines
- Fetch pages of higher quality first
- Continuous operation: get fresh version of already crawled pages

# RESOURCES

- Chapter 19, 20, 21 of  
Introduction to Information Retrieval  
Christopher D. Manning, Prabhakar Raghavan, Hinrich Schütze  
Ebook: <http://nlp.stanford.edu/IR-book/>
- Chapter 5 of  
Mining of Massive Datasets  
Anand Rajaraman, Jure Leskovec, Jeffrey D. Ullman  
Ebook: <http://infolab.stanford.edu/~ullman/mmds.html>